



Design Patterns

Used at Different Software Layers



**Load Balancer
Layer**



**Application
Server Layer**



**API Gateway
Layer**



DesignGurus.io

Client Layer



Client-side Load Balancing: Distributes requests from the client directly to the available service instances.


State Pattern: Allows an object to alter its behavior when its internal state changes.

Composite UI Pattern: Composes responses from multiple microservices to render the complete UI.



DesignGurus.io

Load Balancer Layer



Geographical Distribution: Routes traffic based on the geographical location of the client.

Health Checks: Periodically checks the health of the servers and routes traffic only to healthy ones.

Affinity Based Routing: Routes the user's request to the same server for maintaining session persistence.

Least Connections: Routes traffic to the server with the fewest active connections.

API Gateway Layer



Backend for Frontend (BFF): Tailors API responses to the needs of individual client types.

Circuit Breaker: Detects failures and prevents applications from trying to perform actions that are doomed to fail.

Retry Pattern: Helps to handle transient failures when it tries to connect to a remote service or network resource.

Request Collapsing: Collapses multiple requests for the same operation into a single request.

Web Server Layer



Page Cache Pattern: Stores the output of expensive operations and reuse it to avoid duplicated work.

Compression Pattern: Reduces the size of the response to improve load times.

Lazy Loading: Defers initialization of an object until the point at which it is needed.

Content Negotiation Pattern: The server generates different versions of a resource and serves the one matching the client's criteria.

Application Server Layer



Saga Pattern: Manages long-running transactions and deals with failures and compensating transactions.

CQRS (Command Query Responsibility Segregation): Separates read and write operations to improve performance and scalability.

Proxy Pattern: Provides a surrogate or placeholder for another object to control access to it.

Chain of Responsibility: Passes the request along a chain of handlers.

Caching Layer



Sidecar Caching: Deploy a dedicated cache alongside each microservice to provide isolated and scalable caching functionality.

Cache Chaining: Arrange multiple cache layers hierarchically to handle different granularity or lifetime, querying each layer sequentially on a cache miss.

Time-to-Live (TTL) Caching: Assigns a predefined lifespan to each cache entry, removing or refreshing the entry once its lifespan expires.

CDN Layer



Prefetching: Anticipates user actions and loads resources ahead of time.


Parallel Requesting: Makes multiple requests in parallel to improve load times.

Edge Computing: Processes data closer to the location where it is needed.

Domain Sharding: Splits resources across multiple domains to increase parallel downloads.

Adaptive Image Delivery: Delivers images tailored to the device and user context.

Database Layer



Sharding Pattern: Distributes data across multiple databases to improve scalability.

Replication Pattern: Keeps copies of data in multiple locations for availability and reliability.

Read-Replica Pattern: Uses read replicas to offload read operations from the primary database instance.

Query Object Pattern: An object that represents a database query.